




MATCH THE MODEL TO THE JOB

100x price gap in one vendor's AI range

OpenAI's top model costs 100 times its cheapest per token. Here is how we sort work into three tiers so you stop paying top price for routine jobs.

[SWIPE >>](#)



01 / 06

Sort your work into three piles first

Thinking (strategy, tricky client replies, anything with money or law in it), routine (drafting, summarising, tidying notes) and bulk (sorting hundreds of emails, tagging, extracting). Each pile deserves a different model.

START WITH

List ten things you asked AI to do last week and tag each one.



02 / 06

Top tier is for thinking only

Claude Fable 5.1 and GPT-6 Astra both list at \$10 in and \$50 out per million tokens.

Use them where a wrong answer costs you real money or a client. For a first draft of a thank-you email, it is overkill.

START WITH

Pick one job a month that genuinely needs the top tier.



03 / 06

Mid tier runs the everyday work

Claude Opus 5.5 (\$4/\$20), Claude Sonnet 5.5 (\$2/\$10), GPT-6 Sol (\$2/\$10) and Gemini 3.1 Pro (\$2/\$12) handle most drafting and summarising. Make this your default and move up only when it fails.

START WITH

Set a mid-tier model as your default in your AI app's model picker.



04 / 06

Cheap tier is for bulk jobs

Claude Haiku 4.5 (\$1/\$5), Gemini 3.1 Flash-Lite (\$0.25/\$1.50) and GPT-6 Luna (\$0.10/\$0.50) suit sorting, tagging and extracting at volume. Batch pricing takes roughly another half off on some of them.

START WITH

Run your next bulk job on a cheap tier and check ten results by hand.



05 / 06

Work out a real example before you decide

Say 1,000 customer emails, about 700 tokens in and 200 out each. At list price that is roughly \$17 on Astra, \$3.40 on Sol and \$0.17 on Luna. An illustration, not a promise: your volumes and prompts will differ.

START WITH

Multiply your own monthly volume by the per-million prices above.



06 / 06

Cheaper can mean more retries

A smaller model may miss nuance, need a clearer prompt or take a second go.

Those retries cost tokens and your time too. Test any cheaper tier on ten real jobs before you move a whole workflow onto it.

START WITH

Run the same ten jobs on two tiers and compare the answers side by side.



Pay for thinking only when you think.

We sort every client workflow into thinking, routine and bulk before we pick a model, so the top tier only gets the work that needs it. Prices are USD list prices, checked 5 October 2026.

Which job are you running on a top-tier model that a cheaper one could handle?